Room Climate Control on BuildSim

D7065E, Architecture Document

David Sämfors & Gustav Juul

Version 1.0, 2026-10-12

Contents

1. Summary **2**

2. Use case and building context **2**

3. Requirements **3**

4. Architecture **4**
4.1 Context diagram (C4 Level 1) 4
4.2 Container diagram (C4 Level 2) 5
4.3 Component diagram (C4 Level 3) 6
4.4 Design decisions and alternatives 9

5. Interfaces and data contracts **9**

6. Simulating the sensor values (the physical model) **11**

7. Autonomous services and data pipeline **12**
7.1 The autonomous service 12
7.2 The data pipeline 13

8. Behaviour **13**
8.1 Key scenario (dynamic diagram) 13
8.2 Process lifecycle (state diagram) 14

9. Deployment and component view **14**

10. Test plan **17**

11. Evaluation and results **17**
11.1 Temperature and energy (FR-1, FR-2) 18
11.2 CO_2 and occupancy (FR-3, FR-4) 20
11.3 Cadence and latency (NFR-1) 20
11.4 Fault injection (NFR-2 to NFR-5) 21
11.5 Scaling and the breaking point (NFR-6) 22
11.6 Results against the requirements 23

12. Dashboard **23**

13. Risks and critical reflection **24**

1. Summary

The system heats and ventilates a set of lecture rooms in the LTU A-building, modelled in BuildSim, using one temperature loop and one CO_2 loop per room. People come from the course occupancy simulator; this project contributes the room physics (a two-node heat balance and a CO_2 mass balance), emulated sensor and actuator devices, a decision maker with fault handling, and a monitoring dashboard. Five adjacent lecture rooms (A107–A111, 61–94 seats each) are controlled; which rooms are controlled is one flag per room in a configuration file. The guiding principle is that **a device is a container, but a computation is not**: everything that would be a separate physical device in a real building (each sensor, each actuator, each occupancy counter) runs as its own process so it can fail and be tested independently, while the computations that serve every room (the physics and the decisions) run once for the whole building. That principle came out of a failure, and it threads through Sections 4, 9 and 11.

Field	Value
Project title	Room Climate Control on BuildSim
Team members	David Sämfors, Gustav Juul
Use case	Thermal comfort and indoor air quality in lecture rooms, with a night setback for energy
Repository	https://github.com/Davalav/D7065E
Version / date	v1.0 / 2026-10-12

Changes since the proposal. The design changed five times, each time because something we built or measured forced it:

1. **Monolith to containers and MQTT.** The first version was a set of Python scripts, with room data in a SQLite file. It could not show a device failing or be restarted piece by piece, so every role became its own container talking over an MQTT broker (Section 4.4, D1).
2. **Control moved out of the simulator.** The first script simply moved each room's temperature 20 % of the way to the setpoint per step, and the later simulator computed the heating an ideal thermostat would need. Both hid the controller inside the plant, so nothing about control could be tested. The simulator now only applies whatever power the HVAC actuator holds, and a separate decision maker computes it (D7).
3. **Placeholder decision maker replaced by PID.** Once the containers talked, the decision maker sent a fixed 300 W to prove the message path. It became a clamped PID with anti-windup per loop (D8), with a 21/17 °C day/night schedule for energy and the fault handling of NFR-2 and NFR-3.
4. **Timetable occupancy replaced by the course occupancy simulator.** Occupancy first came from a lecture timetable we converted from Excel to `Lectures.json`. The course's agent-based simulator gives realistic arrivals, lunch dips and corridor flows, and is the same source the 3D viewer shows (D10).
5. **Per-room services consolidated.** Generating eight containers per room for all 119 configured rooms gave 952 containers, and the stack never came up; Section 11.5 shows it would have needed about 23 GB. The decision maker and the simulator are now one container each for all rooms, rooms are enabled by a flag, and five are enabled. The same section measures what this bought, and where the new design breaks (about 55 rooms).

2. Use case and building context

The building. BuildSim models the LTU A-building: three floors, of which level 0 alone has 277 rooms and 45 corridors. It is the *shared state* of the system. It owns the floor plans, keeps a registry of equipment with their sensors and actuators, exposes it all over a REST API, and streams

changes to its 3D viewer over WebSocket. BuildSim provides rooms but no physics: a sensor value only changes when some process writes it. BuildSim is not modified; it is instrumented.

The people. The course occupancy simulator generates a population for each simulated day (we configure 100 people, students, lecturers and staff), books lectures into lecture rooms, moves every person along BuildSim's walkable graph, and publishes positions and per-room counts back to BuildSim. It also owns the simulated clock (60 simulated seconds per real second). Our configuration marks almost all rooms on levels 1 and 2 as unused, so 101 of the 103 simulated people (including three night guards) work on level 0, where the controlled rooms are. Level 0 then has 74 offices, nine lecture rooms and one fika room; all five controlled rooms are lecture rooms.

How the system uses BuildSim.

- The **occupancy publisher** of each room copies that room's head-count from the occupancy simulator into a BuildSim occupancy sensor.
- The **physical simulator** reads each room's HVAC power, ventilation power and occupancy from BuildSim, advances the room physics by one 5-minute step, and hands the new temperature and CO_2 to that room's sensors.
- Each **sensor process** writes its reading into its own BuildSim sensor (so the 3D viewer shows it) and forwards it to the decision maker.
- Each **actuator process** applies commands to its BuildSim actuator and confirms the value BuildSim actually holds. The simulator reads that value next step, which is what closes the loop.
- The **metrics exporter** reads the whole equipment registry for the dashboard.

Actors. The **building manager** watches the dashboard and decides which rooms are controlled; **occupants** only matter through the heat and CO_2 they produce. Neither is on the control path.

The scenario. On a weekday morning the five rooms start cold after the night setback. At 06:00 the setpoint rises to 21 °C and the HVAC loops heat the rooms. From 08:15 lectures of 20–60 students fill some of the rooms: body heat helps the heating, and exhaled CO_2 climbs. When a room passes 900 ppm the ventilation loop opens the damper in proportion, pulling CO_2 back down, then closes it when the lecture ends. At 18:00 the setpoint drops to 17 °C and the heating switches off until the rooms cool to it. Section 11 shows a measured run of this day.

Why it is not trivial. Three things pull against each other. Comfort wants heat early, energy wants it late and low, and the room's thermal mass means heat applied now shows up hours later. CO_2 wants ventilation, but ventilation is energy. And the decision maker only ever sees noisy, possibly stale readings delivered over a network, so it must recognise bad data instead of acting on it.

3. Requirements

Each requirement is a testable statement with a stable ID, traced to the design element that satisfies it and the evidence that verifies it (Section 10). Times are in *simulated* time unless marked "real"; at the configured factor of 60, one simulated 5-minute step lasts 5 real seconds.

ID	Requirement	Prio.	Acceptance criterion	Verified by
FR-1	Regulate each controlled room to its scheduled setpoint	Must	From 15 °C at −10 °C outside, reach 21 °C within 24 h and then stay within ±1 °C; in the live run, ≥ 90 % of daytime steps within ±1 °C once reached	T-1 (unit), T-8 (E1)

ID	Requirement	Prio.	Acceptance criterion	Verified by
FR-2	Night setback: 21 °C from 06:00 to 18:00, 17 °C otherwise	Should	Setpoint changes within one step of 06:00 and 18:00; lower HVAC energy than a constant 21 °C over a day	T-8 (E1), T-9
FR-3	Ventilate on CO_2	Must	Ventilation rises above 0 W whenever CO_2 exceeds 900 ppm; a full room recovers from 1100 to below 1000 ppm within 10 min	T-2 (unit), T-8 (E1)
FR-4	Occupancy drives the physics	Must	Each controlled room's head-count reaches BuildSim every step and enters the heat and CO_2 terms	T-3 (unit), T-8 (E1)
FR-5	Choose controlled rooms without code changes	Should	Setting `"simulate"` in `rooms.json` and regenerating changes the deployed services and nothing else	T-7
NFR-1	Control cadence and latency	Must	Every controlled room gets a temperature and a CO_2 decision every step; reading published → actuator confirmed in under 1 s real at the 95th percentile	T-10
NFR-2	Fail safe on bad sensor data	Must	No reading for 15 s real, or 5 identical readings, puts that loop in fail-safe (heating 0 W, ventilation 600 W) within 1 s real; control resumes on the first good reading	T-4 (unit), T-11
NFR-3	Detect actuator failure	Must	An unconfirmed or wrongly confirmed command is detected within 10 s real and resent until confirmed	T-5 (unit), T-12
NFR-4	Isolate room faults	Must	A failing room's devices do not stop simulation or control of the other rooms	T-6 (unit), T-11, T-12
NFR-5	Recover shared services	Should	Control resumes within 60 s real after the decision maker is killed or the broker is restarted, without manual action	T-13, T-14
NFR-6	Fit on a student laptop	Must	The five-room system runs in under 2 GB RAM, and at least 20 rooms fit in an 8 GB Docker VM	T-15

The tension the autonomous part must resolve is between FR-1 and FR-2 (comfort versus energy, with a slow plant in between) and between FR-3 and energy. The safety requirements NFR-2 and NFR-3 exist because every input arrives over a network: the decision maker must never act on a reading it has reason to distrust. NFR-6 is the requirement the first containerised design broke.

4. Architecture

4.1 Context diagram (C4 Level 1)

The system boundary is the Compose project in our repository. Two course-provided systems sit outside it: BuildSim, which holds the building state and the device registry, and the occupancy simulator, which owns the people and the simulated clock. The system talks to both only over REST. The building manager watches our dashboard and BuildSim's 3D viewer and chooses the controlled rooms by editing `rooms.json`. Occupants are not connected to anything; the occupancy simulator models them and our physics turns their presence into heat and CO_2.

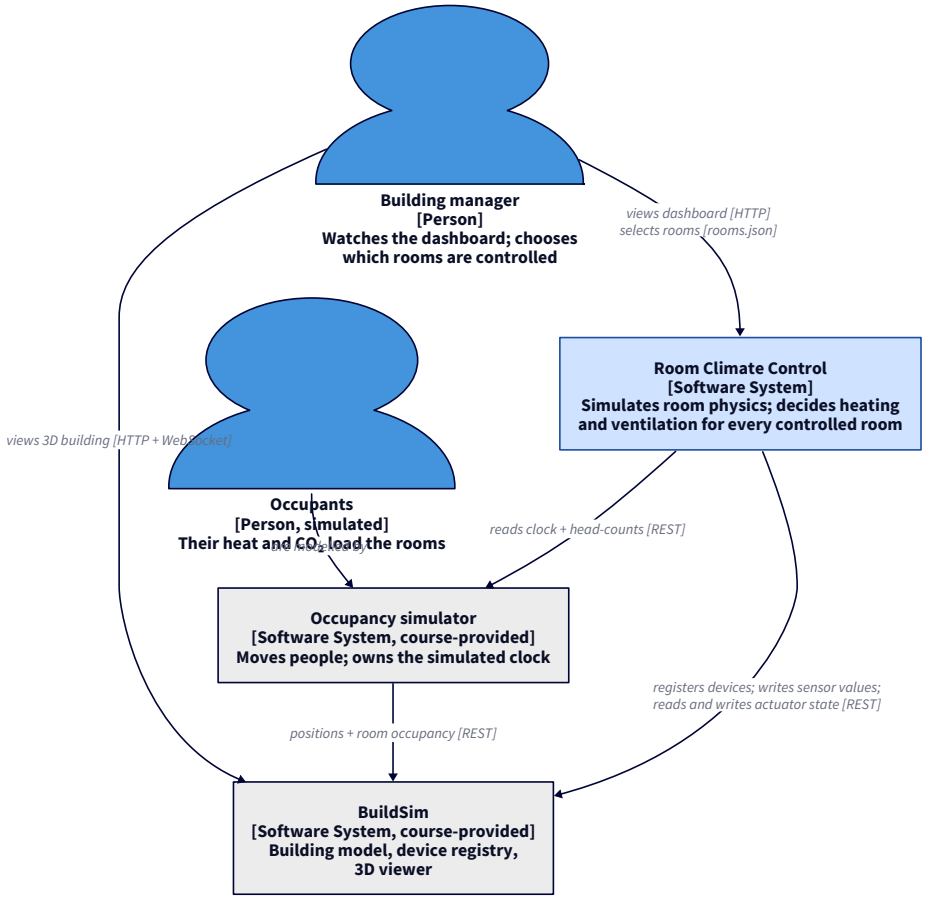


Figure 1: C4 System Context. The system in scope is blue; course-provided systems are grey. Every arrow is labelled with its purpose and protocol.

4.2 Container diagram (C4 Level 2)

Each box is one Docker service in `compose.yaml` or the generated `services-compose.yaml`, and one Python file in the repository. The **per-room device containers** are generated five times per controlled room: an occupancy publisher, a temperature sensor, a CO_2 sensor, an HVAC actuator and a ventilation actuator. The **shared services** run once: the physical simulator, the decision maker, the MQTT broker, the metrics exporter, Prometheus and Grafana. With five rooms that is $25 + 6 = 31$ containers in our project plus the two course systems.

The loop runs clockwise through the figure. The simulator waits for the occupancy simulator's clock to advance five minutes, reads every room's actuator state from BuildSim, steps the physics and publishes each room's physical values to the broker. Each sensor writes its value into BuildSim and forwards it to the decision maker. The decision maker publishes a command per loop; the actuators apply it to BuildSim and acknowledge the value they read back. The next simulator step reads that new actuator state.

Why this split. Each container has one job, and the "reason for separation" column is the part we had to learn the hard way.

Container	Responsibility (one job)	Count	Reason for separation
`occ-*` (`con-occ.py`)	Copy one room's head-count into BuildSim	per room	Emulates a door counter; can fail alone

Container	Responsibility (one job)	Count	Reason for separation
`sensor-*`, `sensor-co2-*` (`con-sensor.py`)	Own one BuildSim sensor; forward readings	2 per room	A real sensor fails independently; fault tests stop exactly one
`hvac-*`, `vent-*` (`con-hvac.py`, `con-vent.py`)	Apply commands to one BuildSim actuator; acknowledge	2 per room	Separate failure domains; the ack proves the value was applied
`sim` (`con-sim.py`)	Room physics for every room	1	The environment, not a device: one clock and one step for all rooms
`dec` (`con-dec.py`)	Two PID loops per room, fault detection	1	Decision logic evolves independently of devices; one process holds all controller state
`mosquitto`	Route readings, commands, acks	1	Decouples devices from the decision maker
`metrics-exporter`	Expose BuildSim state as Prometheus metrics	1	Monitoring must not load the control path
`prometheus`, `grafana`	Store recent metrics; dashboard	1 each	Off-the-shelf; replaceable without touching control

4.3 Component diagram (C4 Level 3)

The decision maker is the one container with internal structure worth drawing, because since the consolidation it serves every room. At start-up the **controller registry** reads `rooms.json` and builds two independent controllers per simulated room. The **topic router** subscribes to each controller's reading and acknowledgement topics (20 subscriptions for five rooms) and hands every message to exactly one controller, so rooms share a process but no state. Inside a controller, the **sensor health monitor** decides whether a reading can be trusted; trusted readings go to the **PID regulator** with the setpoint from the **setpoint scheduler**, distrusted ones trigger the fail-safe output. The **command publisher** numbers every command, and the **actuator ack monitor** expects that number back. A **health watchdog** thread checks every controller once a real second, so a sensor that goes silent is caught even though no message arrives to trigger the check.

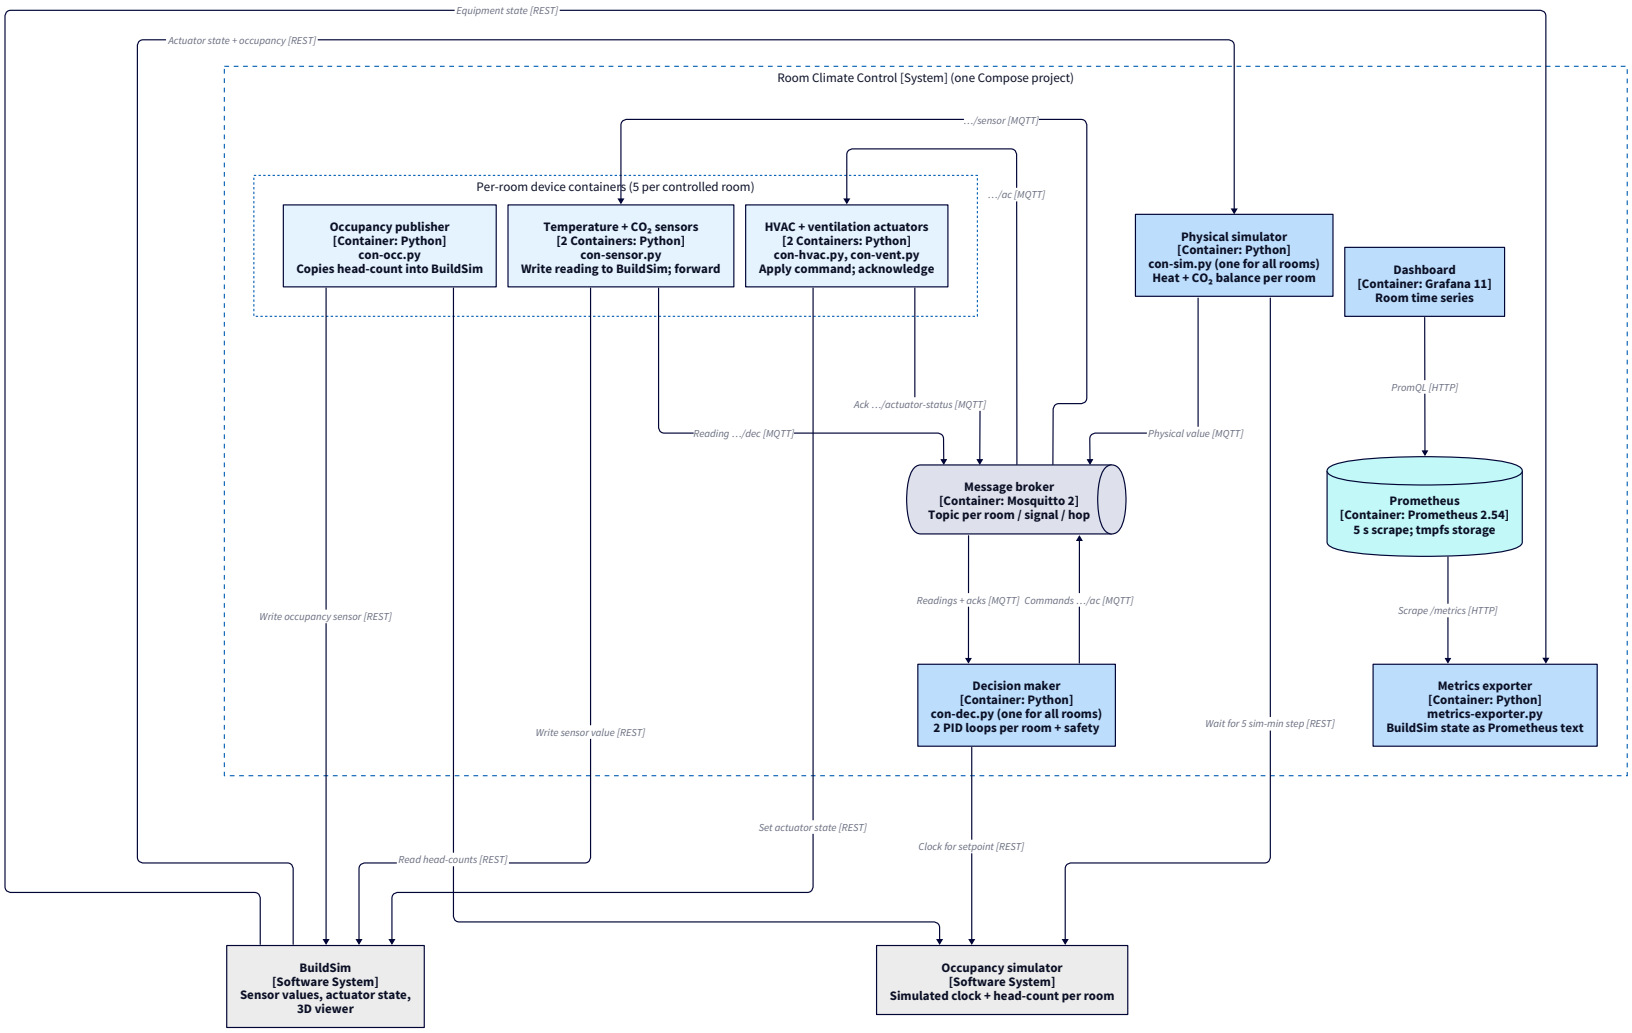


Figure 3: C4 Container diagram. The dashed boundary is the system (one Compose project); the inner dotted group is instantiated once per controlled room. External systems are grey.

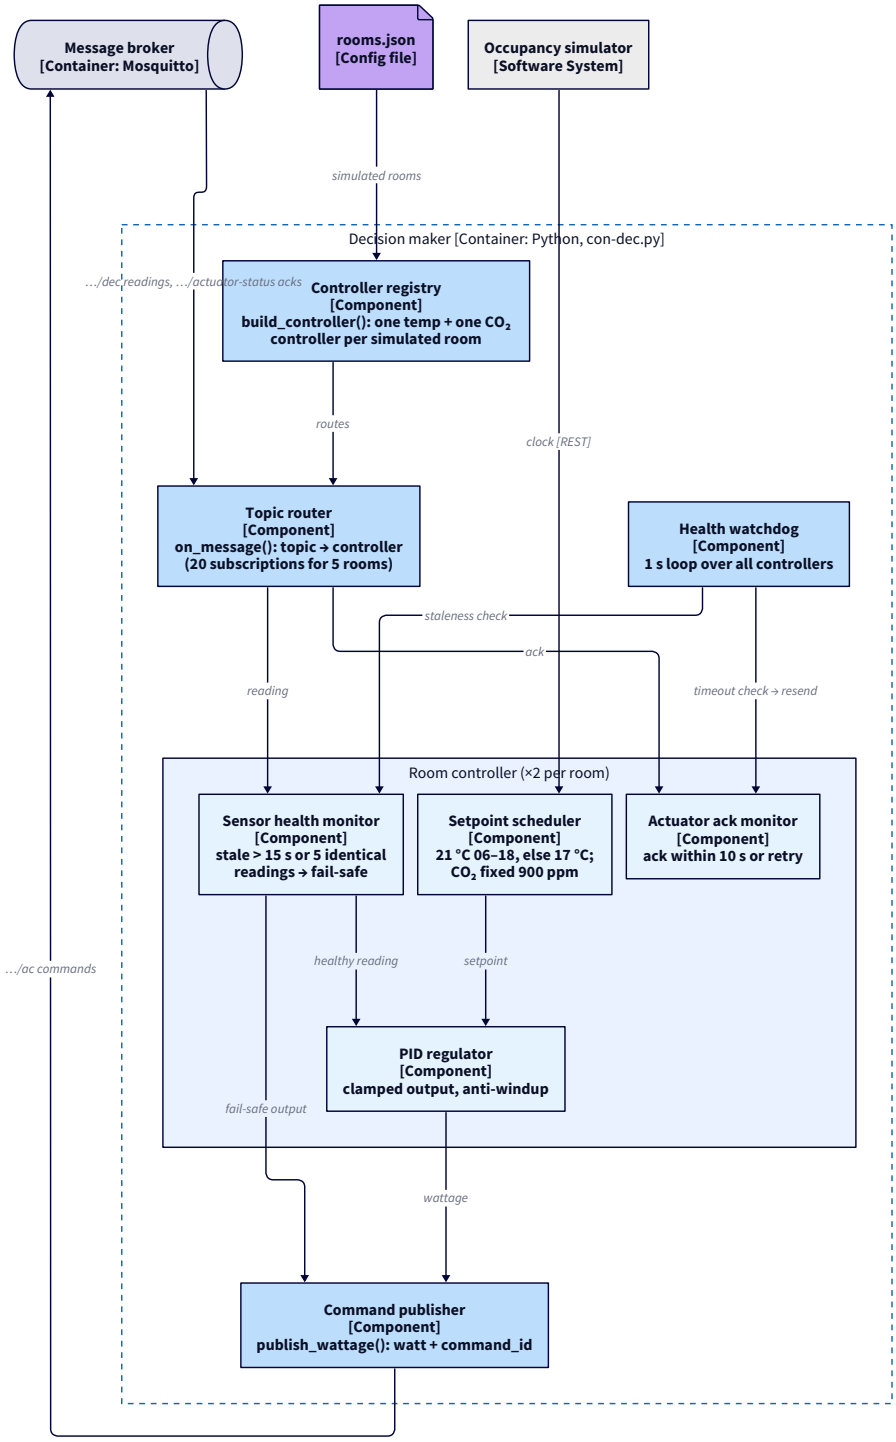


Figure 2: C4 Component diagram of the decision maker. The inner box is instantiated twice per room (temperature and CO_2).

4.4 Design decisions and alternatives

	Decision	Alternatives considered	Why this one (requirement)	Trade-off accepted
D1	One container per role, MQTT between them	Monolith with a shared SQLite file	Device faults can be injected and recovered one at a time (NFR-2, NFR-4); roles restart independently (NFR-5)	More processes and a broker to run
D2	One decision maker for all rooms	One per room and signal (2 per room); one per room	952 containers for 119 rooms never started and would need about 23 GB (NFR-6, Section 11.5); routing by topic keeps rooms isolated in one process (NFR-4)	One process is a single point of failure for control; mitigated by restart and fail-safe (NFR-5)
D3	One physical simulator for all rooms	One simulator per room	The physics is the environment, not a device: one clock read and one step per timestep; a failing room is skipped, not fatal (NFR-4)	A simulator crash stops physics for all rooms until restart
D4	Keep sensors and actuators per room	Merge them into one device container	They are what fails in a real building, and per-device fault tests need them separate (NFR-2, NFR-3)	Memory grows by five containers per room (Section 11.5)
D5	MQTT pub/sub for readings, commands and acks; topic per room, signal and hop	REST calls between services; one topic per room with a type field	Producers do not need to know consumers; the router maps a topic straight to a controller (NFR-1)	QoS 0, no persistence: a message sent during a broker outage is lost (Section 11.4)
D6	Sensor process writes BuildSim; simulator hands values to it over MQTT	Simulator writes BuildSim and sensors poll it (course reference flow)	Each BuildSim sensor has exactly one writer, its device, and a dead sensor visibly stops updating BuildSim (NFR-2)	The "physical" link is modelled as a network message; one extra hop
D7	Plant and controller separated	Ideal thermostat inside the simulator	Control can only be tested if the plant just applies actuator power (FR-1)	Must tune a controller instead of getting perfect control for free
D8	Clamped PID with anti-windup	On/off with hysteresis; MPC or learned policy	Proportional power with no oscillation, three parameters, and a closed-loop unit test that reaches and holds the setpoint (FR-1, FR-3)	Gains are tuned against real seconds, so they depend on the clock factor (Section 13)
D9	Fail-safe: heating 0 W, ventilation 600 W	Hold last command	A frozen value must not keep driving the actuator; fresh air beats energy when CO_2 is unknown (NFR-2)	Rooms cool during a sensor outage
D10	Course occupancy simulator	Own timetable from `Lectures.json`; random density sampler	Realistic agent movement and lectures, same data as the 3D view (FR-4)	External dependency with a broken health check (Section 9)
D11	Time-scheduled setback (21/17 °C)	Occupancy-based setback	Predictable and needs no occupancy forecast; lecture occupancy is bursty (FR-2)	Empty rooms are heated all day (Section 13)
D12	Prometheus + Grafana, dashboard as code	Custom web UI or BuildSim session overlay; time-series DB	Provisioned with the stack, zero code for plots (Section 12)	No history across restarts; decisions not shown
D13	Generate per-room services from `rooms.json`	Hand-written compose file	One flag per room (FR-5); the 400 KB generated file was unmaintainable by hand	Regenerate and rebuild after every change

5. Interfaces and data contracts

Two patterns are used, each where it fits. **MQTT** (Mosquitto 2, port 1883, QoS 0, not retained) carries everything that flows between our own services, because a reading has a producer and a

consumer that should not know each other, and because the decision maker can subscribe to every room on one connection. **REST** (HTTP + JSON) is used for every call to BuildSim and the occupancy simulator: each is a one-to-one request whose result the caller needs, and both systems only offer REST. All payloads are JSON objects; numeric values are sent as strings, matching BuildSim's `"data_type":"text"` sensors.

MQTT topics. Every topic has the form `{level}/{room}/{signal}/{hop}`, where `level` is the floor number (`0`), `room` the URL-quoted BuildSim room name, `signal` is `temp` or `co2`, and `hop` says where the message is in the loop. One topic per hop makes each process's subscriptions explicit and lets the decision maker map a topic to a controller with a dictionary lookup.

Topic	From → to	Payload (example)	Rate
`0/A108/temp/sensor`	sim → temp sensor	`{"value":"20.04"}` (°C)	1 per step
`0/A108/co2/sensor`	sim → CO_2 sensor	`{"value":"937.6"}` (ppm)	1 per step
`0/A108/temp/dec`	sensor → dec	`{"value":"20.04"}`	1 per step
`0/A108/co2/dec`	sensor → dec	`{"value":"937.6"}`	1 per step
`0/A108/temp/ac`	dec → HVAC	`{"watt":"6345.2","command_id":41}` or `{"set_temp":"21"}`	1 per step; setpoint on change
`0/A108/co2/ac`	dec → vent	`{"watt":"188.0","command_id":41}` (0–600)	1 per step
`0/A108/temp/actuator-status`	HVAC → dec	`{"command_id":41, "state":"6345.2"}`	per command
`0/A108/co2/actuator-status`	vent → dec	`{"command_id":41, "state":"188.0"}`	per command

Command semantics. `command_id` increases by one per command and per loop. An acknowledgement is accepted only if its `command_id` equals the latest command's and its `state` equals the commanded value within 0.01; anything else counts as a failed apply. An unacknowledged command is resent unchanged after 10 s real (Section 7.1). The `set_temp` message is informational: the HVAC actuator stores it in BuildSim for the viewer and the dashboard, but the power command is what heats.

BuildSim REST calls. Device ids follow fixed patterns, so any process can compute another's ids from the room name.

Caller	Call	Body / result	When
sensor, HVAC, vent, occ	`POST /api/equipment`, then `POST /api/equipment/{eq}/sensors` or `/actuators`	equipment `temp-A108`, `co2-A108`, `hvac-A108`, `vent-A108`, `occupancy-A108`; devices `A108-temp`, `A108-co2`, `A108-watt`, `A108-set`, `A108-vent`, `A108-occupancy`	start-up
sensor, occ	`PUT /api/sensors/{id}/value`	`{"data_type":"text", "value":"20.04"}`	each reading
HVAC, vent	`PUT /api/actuators/{id}/state`, then `GET /api/actuators/{id}`	`{"state":"6345.2"}`; the GET's `state` is acknowledged	each command
sim	`GET /api/actuators/A108-watt`, `.../A108-vent`, `GET /api/sensors/A108-occupancy`	`state`, `value` as strings	each step, per room
sim	`GET /api/building/floors/level0`	room `area`	start-up
exporter	`GET /api/equipment`	every device with its current value	each scrape (5 s)

Occupancy simulator. `GET /api/state` returns, among other things, the clock and a list of rooms; the simulator and occupancy publishers use these two fields:

```
{"sim": {"clock": {"time": "2026-10-12T08:24:53Z", "time_of_day": "08:24",
                   "weekday": "Monday", "factor": 60, "running": true},
         "rooms": [{"name": "A108", "level": "level0", "role": "lecture",
                    "area_m2": 122.875, "capacity": 61, "occupants": 29}]}}
```


Monitoring. `GET /metrics` on the exporter returns Prometheus text with eight gauges: `room_temperature_celsius`, `room_co2_ppm`, `room_occupancy_people`, `room_power_consumption_watts`, `room_ventilation_command_watts` and `room_set_temperature_celsius`, labelled `room`, plus `building_outside_temperature_celsius` and `building_occupancy_sim_time_seconds`.

6. Simulating the sensor values (the physical model)

BuildSim has rooms but no physics, so the simulator (`con-sim.py`) produces every temperature and CO_2 value. It is deliberately simple, but it has the three properties the controller needs: slow thermal dynamics, a CO_2 response to people, and a genuine response to both actuators.

Time. The simulator polls the occupancy simulator's clock every 0.5 s real and steps when five simulated minutes have passed ($\Delta t = 300$ s). Outside temperature is interpolated linearly from `Weather.json`, which has one value every three hours for Monday to Friday.

Occupancy. The head-count n of each room comes from the occupancy simulator via the room's occupancy sensor in BuildSim. The occupancy simulator is agent-based: each person gets a timetable for the day (arrival 08:00–10:00, optional fika breaks, lunch 12:00–13:00 in or out of the building, departure 16:00–19:00); students attend one to three lectures in five slots between 08:15 and 18:00, and lectures are packed into the fewest rooms that seat the demand. People walk between rooms along BuildSim's walkable graph at 1.3 m/s. All five of our rooms classify as lecture rooms (61–94 seats), so their occupancy jumps from 0 to 20–60 at slot boundaries, a hard test for the CO_2 loop.

Temperature. Each room is two thermal nodes: the air (temperature T, capacity $C_a = \rho c_p V$) and the room's thermal mass (temperature T_m, capacity $C_m = c_m A$):

$$C_a \frac{dT}{dt} = P_{\text{hvac}} + q_{\text{occ}}\, n - (U A_w + \rho c_p V \tfrac{\text{ACH}_0}{3600})(T - T_{\text{out}}) + hA(T_m - T), \tag{1}$$

$$C_m \frac{dT_m}{dt} = -hA(T_m - T), \tag{2}$$

with floor area A, volume $V = 3A$, exterior wall area $A_w = 3\,L_w$ from the wall length L_w in `rooms.json`, $U = 0.25$ W/m²K, background air change $\text{ACH}_0 = 0.5$ h^{-1}, $c_m = 165$ kJ/m²K, air–mass coupling $h = 8$ W/m²K, and $q_{\text{occ}} = 75$ W sensible heat per person. P_{hvac} is the HVAC actuator's power as read from BuildSim. The equations are integrated with explicit Euler in sub-steps of at most 60 s (six 50 s sub-steps per step), which keeps the air node stable.

CO_2. With ventilation power P_v mapped linearly to 0–3 air changes per hour, $\lambda = 3\,\min(P_v/600, 1)/3600$ s^{-1}, and generation $g = n\,G/V \cdot 10^6$ ppm/s with $G = 5$ mL/s per person, the well-mixed mass balance has the exact solution over a step

$$C(t + \Delta t) = C_{\text{eq}} + (C(t) - C_{\text{eq}})e^{-\lambda \Delta t}, \qquad C_{\text{eq}} = 420 + g/\lambda, \tag{3}$$

and $C(t + \Delta t) = C(t) + g\,\Delta t$ when $\lambda = 0$. Using the exact solution rather than an Euler step means a 5-minute step cannot overshoot below the outdoor 420 ppm.

Sensor noise. Each published value gets Gaussian noise ($\sigma = 0.2$ °C, $\sigma = 10$ ppm, CO_2 clamped at 420 ppm). The noise is not cosmetic: it is what lets the stuck-sensor rule (five identical readings) work, because a live sensor never repeats a value exactly.

Closing the loop. Both actuators enter the equations: P_{hvac} directly as heat, P_v through λ. A command applied in step k is read at the start of step $k+1$, so actuation lags one step (5 min), the same as a real BMS polling cycle.

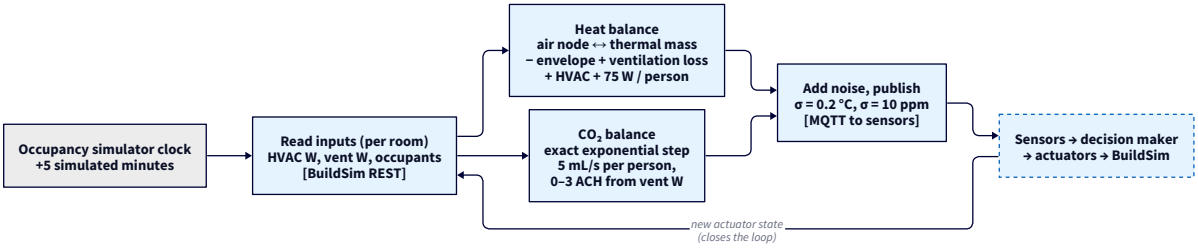


Figure 4: One simulator step. The dashed path is the rest of the system; its effect re-enters as actuator state at the next step.

7. Autonomous services and data pipeline

7.1 The autonomous service

The decision maker (`con-dec.py`) runs two control loops per room: heating power from temperature, and ventilation power from CO_2.

Algorithm. Each loop is a PID regulator acting on every incoming reading:

$$u = \text{clamp}_{[0,\,u_{\max}]}\Big(K_p e + K_i \int e\,dt - K_d \frac{dy}{dt}\Big),$$

with $e = T_{\text{sp}} - T$ for heating and $e = C - C_{\text{sp}}$ for ventilation. The derivative acts on the measurement, not the error, so a setpoint change does not kick the output. The integral is only updated when the unclamped output is inside $[0, u_{\max}]$ or the error would move it back inside (conditional integration), which prevents wind-up during the hours-long warm-up.

Loop	Setpoint	K_p	K_i	K_d	$u_{\max}$	Fail-safe
Heating	21 °C (06–18), 17 °C	1000 W/°C	10 W/(°C·s)	0	30 kW	0 W
Ventilation	900 ppm	5 W/ppm	0.001 W/(ppm·s)	0	600 W	600 W

Why PID. The requirement is to hold a setpoint on a slow, first-order-ish plant with a known actuator range. A PID does that with three interpretable parameters, its behaviour can be verified in a closed-loop unit test against the same physics code (T-1, T-2), and its output is proportional, so it does not cycle the HVAC on and off the way a hysteresis thermostat would. A learned or optimising controller would need a model or data we do not have, and would be harder to verify for a safety-relevant function. The setpoint schedule is a simple clock rule for the same reason (D11).

Inputs, outputs, integration. Inputs are the readings on `.../dec` and the clock from the occupancy simulator (for the setpoint). Outputs are the numbered commands on `.../ac`. All controllers share one MQTT connection; the router delivers each message to one controller, and malformed payloads or an unreachable occupancy simulator are caught and logged per message without touching the other controllers. An unexpected exception type, however, would propagate into the shared MQTT network thread and stop delivery to every room; the handlers catch only the errors we have seen.

Failure and safety behaviour.

- **Stale or stuck input.** The sensor health monitor flags a loop if no reading arrived for 15 s real (three missed steps) or the last five readings were identical. The loop then publishes its fail-safe output once and ignores readings until a fresh, changed one arrives. The 1 s watchdog thread makes the stale check independent of message arrival.
- **Actuator does not apply.** Every command carries an id; if no matching acknowledgement arrives within 10 s real, the watchdog resends the last command and restarts the timer, until it is confirmed.
- **Decision maker unavailable.** Actuators hold their last value. Docker restarts a crashed process (`restart: unless-stopped`); controllers start empty and resume from the next reading. Section 11.4 shows the gap in this: a container stopped with `docker kill` is not restarted.

7.2 The data pipeline

The control path keeps no history on purpose: each controller holds only its last reading, its integral and its pending command, and the readings themselves arrive over MQTT. Storing and querying readings on the control path would add latency and a failure mode for no control benefit, since the PID only needs the current value.

For monitoring, a separate pipeline samples BuildSim, which already holds the latest value of every sensor and actuator. On each Prometheus scrape (every 5 s real, i.e. every simulated 5 minutes) the metrics exporter reads the whole equipment registry and the occupancy clock and returns eight gauges labelled by room. Prometheus stores them in a tmpfs volume, so history lasts only as long as the container: this is a deliberate choice for a laptop demo (nothing to clean up, no disk growth) and a real limitation for long-term analysis (Section 13).

Data quality. Missing and stale readings are handled where they matter, in the decision maker (above). The exporter skips values it cannot parse and logs them, and returns HTTP 503 when BuildSim or the occupancy simulator is unreachable, so Prometheus records a gap instead of a wrong number. Duplicates cannot arise because the exporter reads current state, not a stream.

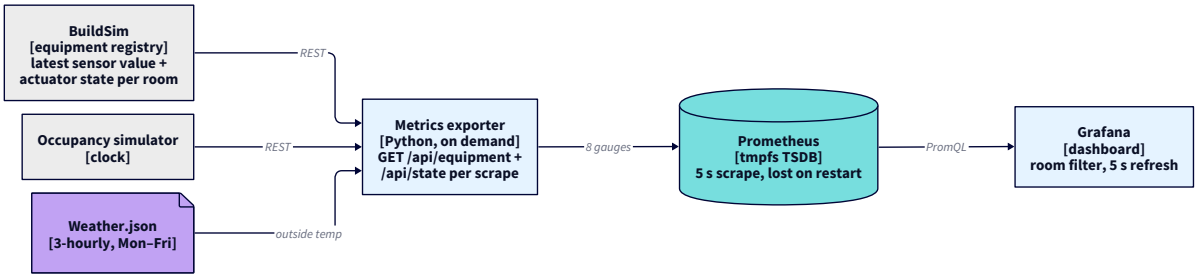


Figure 5: Monitoring pipeline. Control does not read from it (dashed arrow).

8. Behaviour

8.1 Key scenario (dynamic diagram)

The scenario is the start of a lecture, taken from the logs of the measured run (Section 11) with real timestamps and values. At 08:09 on Monday 29 students arrived in A108; CO_2 crossed 900 ppm at 08:14 and the decision maker answered with 32 W of ventilation. Figure 7 follows the next step. The simulator steps when the clock reaches 08:19:47 (1), reads A108's actuators and head-count from BuildSim (2), and computes 956.6 ppm including sensor noise (3). The CO_2 sensor writes the value into BuildSim, where the 3D viewer shows it, and forwards it (4–7). The router hands it to A108's CO_2 controller, whose PID turns 56.6 ppm of excess into 283.32 W (8–9).

The ventilation actuator applies it, reads it back and acknowledges command 7 (10–14): 50 ms after the decision and 93 ms after the simulator published, measured on the container clock. The step at 08:24:53 reads the new power (15): 1.4 air changes per hour pull CO_2 back to 911.6 ppm.

Two timing facts are visible here that static diagrams hide. The actuation lags one step, because the simulator reads actuator state at the start of a step; and the whole message chain takes about a tenth of a second, so the 5 s step, not the network, sets the control rate.

8.2 Process lifecycle (state diagram)

Each of the decision maker's controllers (two per room) has the lifecycle below; every transition is a branch in `con-dec.py` or `control_safety.py`. A controller starts in *Waiting* and has 15 s real to receive its first reading before it fails safe, so a room whose sensor never starts is not left uncontrolled. *Controlling* and *Fail-safe* describe the sensor side; *Awaiting ack* and *Retrying* describe the actuator side and run alongside them. The fault-injection runs in Section 11.4 drive a controller through every transition.

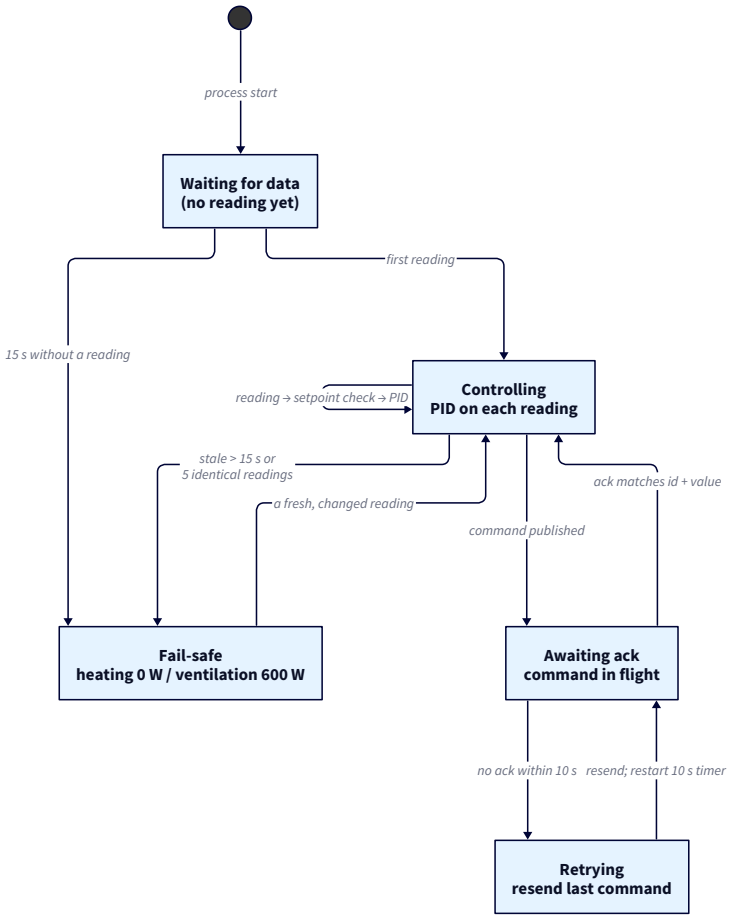


Figure 6: Lifecycle of one controller in the decision maker.

9. Deployment and component view

Everything runs on one laptop: Docker Desktop's Linux VM (8 GB, 12 vCPU on the measurement machine) hosts one Compose project on a private bridge network, where services find each other by name (`buildsim:9090`, `mosquitto:1883`). Only the browser-facing ports are published, and only on 127.0.0.1: Grafana 3000, Prometheus 9091, BuildSim 9090, occupancy UI 8081, MQTT

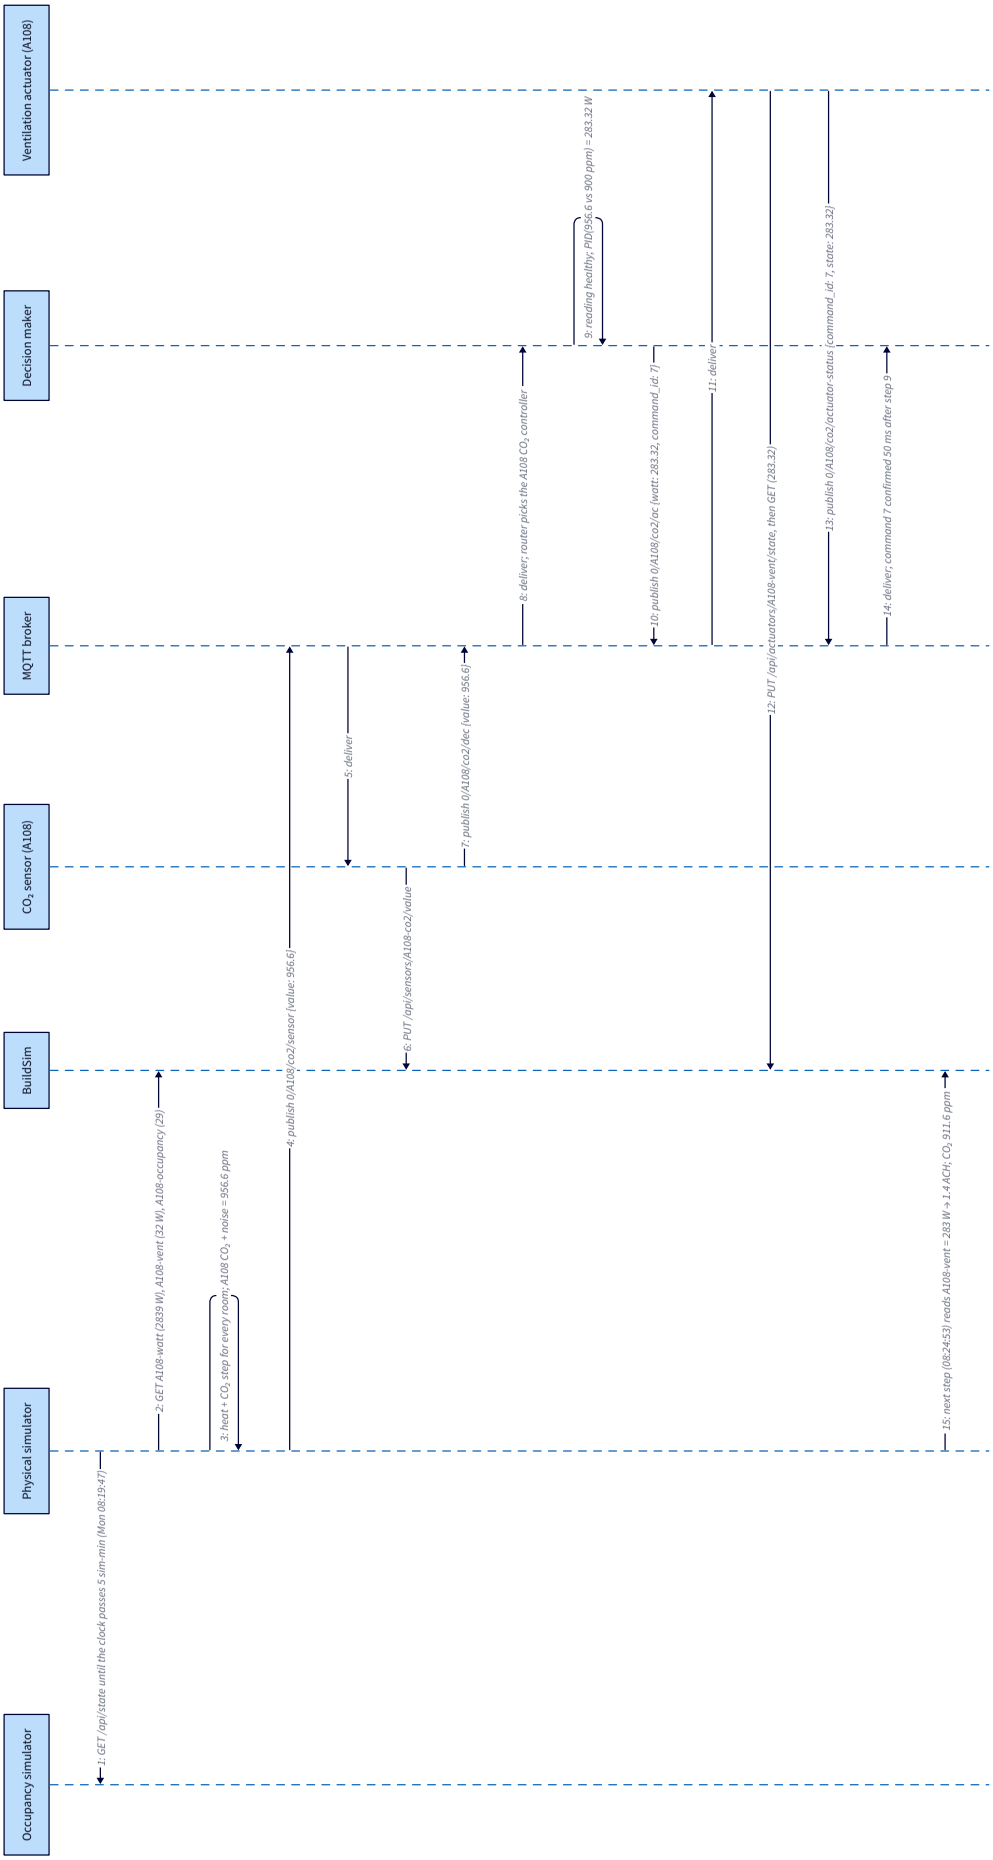


Figure 7: C4 Dynamic diagram: one simulation step for the CO_2 loop of room A108 during a lecture.

1883. All our services are built from one image (`d7065e-python`, Python 3.12-slim with `requests` and `paho-mqtt`); each container runs a different script with different arguments.

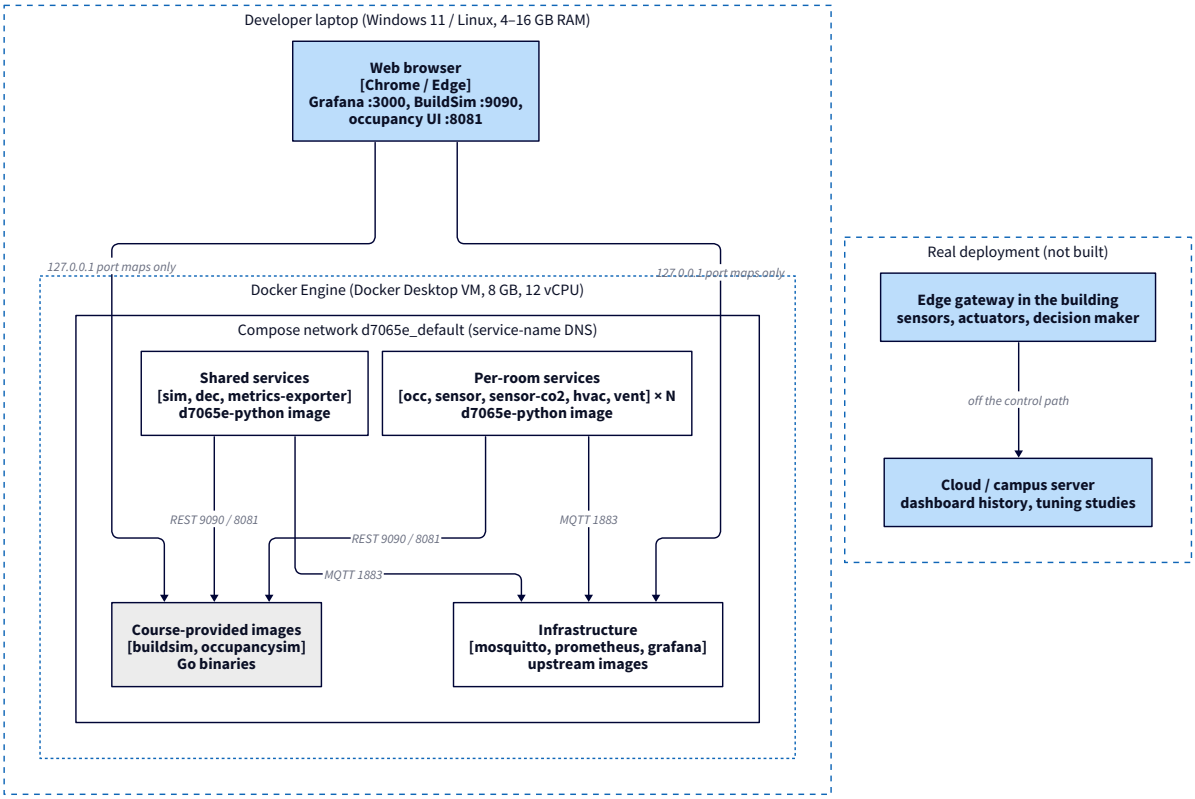


Figure 8: Deployment. Everything in the dashed laptop box is what was built; the lower box is where the parts would go in a real building.

Placement and its consequences.

- **Fault isolation.** Every device is its own container, so stopping one sensor is a real, observable failure, which is how T-11 and T-12 are run. The price is memory: each Python container costs about 19 MiB, or 23 MiB of VM memory once Docker's per-container overhead is counted (Section 11.5).
- **Shared services as single points of failure.** The decision maker and simulator serve every room. Splitting them per room would remove that, but would cost two containers per room and gave us a stack that did not start at 119 rooms. We kept them shared and made them recover instead (restart policy, in-process isolation, NFR-5).
- **Start-up ordering.** Compose starts the course systems first; everything else waits for BuildSim's health check and for the broker. The occupancy simulator's image has no health command (its binary would start a second simulator instead), so its health check is disabled and its dependents retry until its API answers.
- **Edge versus cloud.** In a real building the loop (sensors, actuators, broker, decision maker) belongs on an edge gateway in the building, so heating and ventilation keep working when the network to the outside is down; the dashboard history and any tuning studies could move to a campus or cloud server, off the control path. Nothing in our loop needs external compute: one decision costs microseconds.

10. Test plan

Three levels of evidence are used: **unit tests** that run the real controller and physics code without containers (`python -m unittest discover -s tests`, nine tests, all passing), **end-to-end runs** of the full stack with the measurements of Section 11, and **fault injection** by stopping or killing individual containers during a run. Each row traces to a requirement.

ID	Test	Requirement	Expected result / evidence
T-1	`test_empty_room_reaches_temperature_setpoint_at_minus_10_c`: PID + physics, empty 750 m² room, 15 °C start, −10 °C outside, 48 h	FR-1	21 °C within 24 h; within ±1 °C from 42–48 h
T-2	`test_full_room_vent_recovers_below_1000_ppm_within_10_minutes`: 94 people, 1100 ppm start	FR-3	Below 1000 ppm within 10 min
T-3	Three `test_occupancy_publisher_*` tests	FR-4	Correct room and level selected; missing room and negative counts rejected
T-4	`test_crashed_or_stuck_sensor_enters_failsafe_and_recovers`, both loops	NFR-2	Fail-safe 0 W / 600 W on silence and on five identical readings; recovery on a new value
T-5	`test_stuck_or_crashed_actuator_is_detected_retried_and_recovers`	NFR-3	Resend after 10 s; wrong value rejected; matching ack clears
T-6	`test_single_decision_maker_routes_each_room_to_its_own_controller`, `test_simulator_keeps_other_rooms_running_when_one_room_fails`	NFR-4	Message reaches only its room's controller; a failing room is skipped, the others are published
T-7	Toggle `"simulate"` for rooms, run `yaml_generate.py`	FR-5	5 services per enabled room + `dec` + `sim`; `docker compose config` valid
T-8 (E1)	40 simulated hours, five rooms, full stack	FR-1–FR-4	Temperature, setpoint, power, CO_2, ventilation, occupancy per step (Section 11.1–11.2)
T-9	Offline replay of E1's occupancy and weather with a constant 21 °C setpoint	FR-2	HVAC energy with and without setback (Section 11.2)
T-10	Timestamps in E1's container logs	NFR-1	Latency distribution per hop (Section 11.3)
T-11 (E2)	Stop A110's temperature sensor for 60 s, then start it	NFR-2, NFR-4	Fail-safe within 1 s of the 15 s limit; other rooms unaffected; recovery
T-12 (E3)	Stop A109's HVAC actuator for 40 s, then start it	NFR-3, NFR-4	Retry within 10 s; confirmation after restart
T-13 (E4)	`docker compose kill dec` (SIGKILL)	NFR-5	Control resumes within 60 s without manual action
T-14 (E5)	Stop the MQTT broker for 20 s, then start it	NFR-5	All loops resume within 60 s
T-15 (stress)	`docker stats` at 5 rooms; the full stack at 20 and 40 rooms	NFR-6	Memory per container type; memory versus rooms; breaking point (Section 11.5)

11. Evaluation and results

Everything in this section was measured on the running system, or replayed through its code, on one laptop (Intel i7-1355U, 16 GB; Docker Desktop VM with 8 GB and 12 vCPU). The raw logs and the scripts that turn them into these numbers are in `report/experiments/` and `report/data/`.

Experiment E1. The five-room stack ran for 49 real minutes, 581 simulation steps covering 48 simulated hours from Monday 07:49 to Wednesday 09:02, followed by the fault injections E2–E5. BuildSim had been running before E1, so the rooms started from the state the previous run left:

air at 17.6–18.2 °C and CO_2 already elevated in A108 (858 ppm) and A110 (889 ppm). Every number below comes from Docker log timestamps, which share one clock, so host/VM clock skew does not enter them.

11.1 Temperature and energy (FR-1, FR-2)

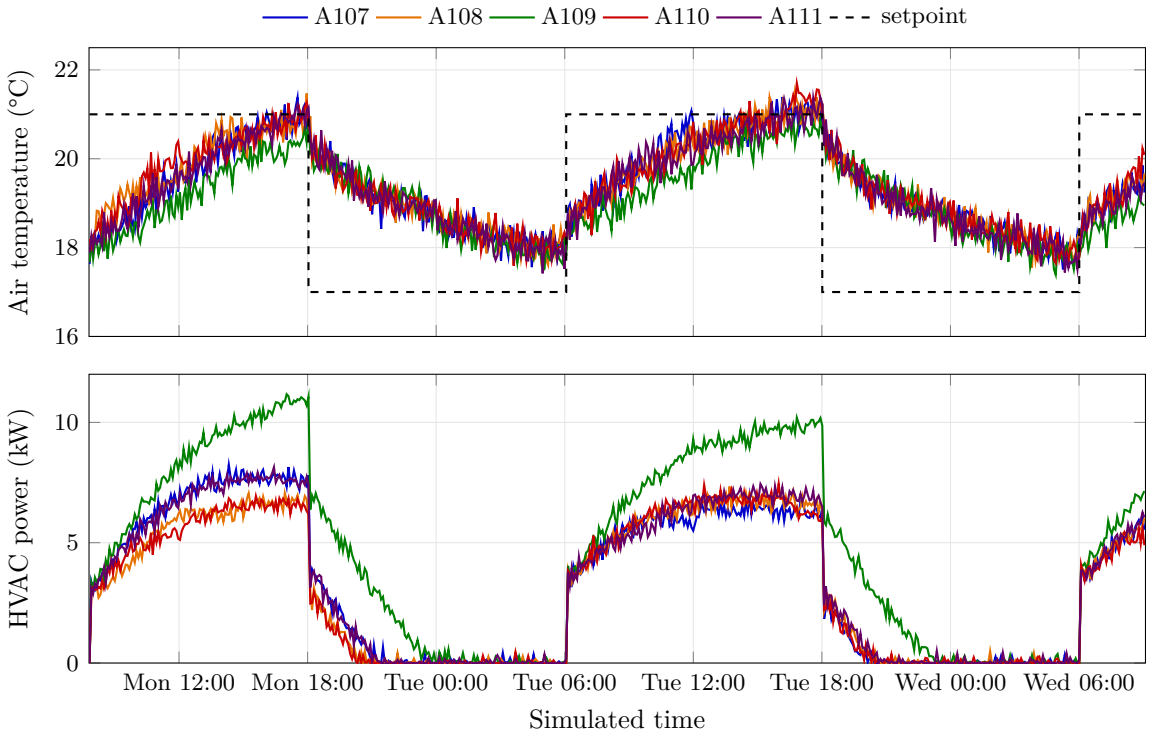


Figure 9: E1: measured air temperature (with sensor noise) and HVAC power per room. The dashed line is the scheduled setpoint.

Figure 9 shows the two simulated days. On Monday the rooms needed 5.4–10.1 hours to climb from 17.6–18.2 °C to 20.5 °C, although the controllers never used more than 11.2 kW of their 30 kW: the integral term builds slowly because it integrates real seconds (Section 13), and the four-times oversized thermal mass absorbs the rest. Once warm, the loops held the band: on Tuesday afternoon the rooms averaged 20.6–21.1 °C. At 18:00 the power halved in the first step, the proportional term reacting to the new setpoint, but took 2.1–6.0 hours to reach zero, because the anti-windup limits the stored integral without resetting it on a setpoint change. Overnight the rooms coasted towards 17 °C with almost no heating: night energy was 6–19 % of day energy. The weak point is the morning: after the 06:00 step the rooms needed 2.4–6.3 hours to get within 1 °C of 21 °C, so lectures starting at 08:15 met rooms at 18–19 °C. Over all occupied steps only 52 % (95 of 181) were within ±1 °C. FR-1 is met by the controller in isolation (T-1) but not by the deployed system.

Room (area[a], wall)	Reached 20.5 °C after	Daytime steps within ±1 °C[b]	Occupied steps within ±1 °C (steps)	Tuesday 06:00 → within 1 °C	Energy day / night (kWh)
A107 (490, 8 m)	5.8 h	63 %	82 % (27)	3.1 h	148 / 10.5
A108 (492, 8 m)	5.4 h	62 %	43 % (47)	3.6 h	141 / 8.8
A109 (752, 16 m)	10.1 h	35 %	83 % (24)	6.3 h	193 / 35.9
A110 (497, 8 m)	6.9 h	59 %	47 % (57)	2.4 h	141 / 7.9
A111 (510, 16 m)	6.2 h	60 %	23 % (26)	3.0 h	151 / 11.6

[a] Area as the simulator uses it (BuildSim plan units², see Section 13); real floor area is a quarter of this. [b] Steps between 06:00 and 18:00, counted from the first time the room reached 20.5 °C.

Night setback (FR-2). In the decision maker's log the setpoint changed in the first step after each of the four 06:00 and 18:00 boundaries in E1 (at simulated 18:02, 06:04, 18:01 and 06:00). The energy effect cannot be read from E1 alone, because there is no run without the setback, so it was replayed: E1's occupancy and outside temperature were fed step by step through the same simulator and PID code, with the same noise and actuation lag (Figure 10). The replay is good enough for this: with the deployed schedule it reproduces the measured room temperatures with an RMSE of 0.21 °C in every room (the sensor noise alone is 0.2 °C) and the measured energy within 2.3 %. Over the 48 hours the schedule used 868 kWh against 1246 kWh for a constant 21 °C, **a saving of 30 %**. The price is comfort: with a constant setpoint the same controller keeps 99–100 % of daytime steps within ±1 °C, against 53–59 % with the setback. The controller can hold the band; the schedule is what loses it, which is why Section 13 proposes an optimal-start setback rather than a better PID. FR-2 is met.

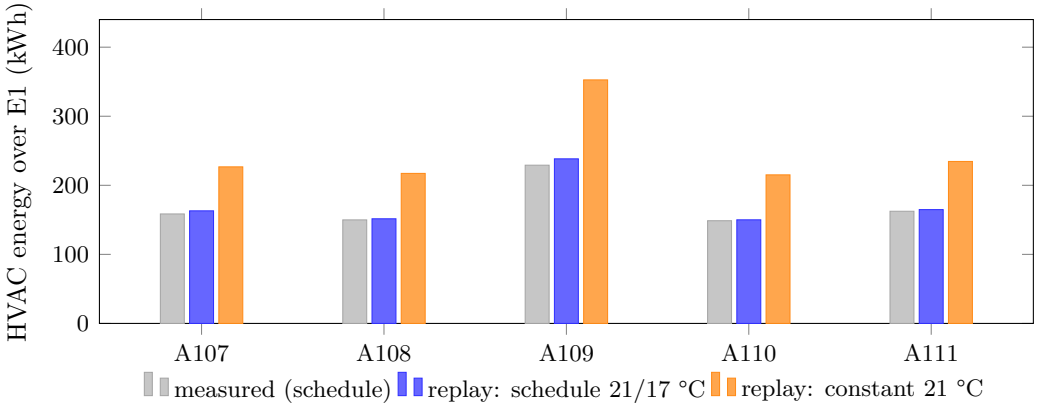


Figure 10: HVAC energy over the E1 window per room. The replay feeds E1's measured occupancy and weather through the production physics and PID code.

11.2 CO_2 and occupancy (FR-3, FR-4)

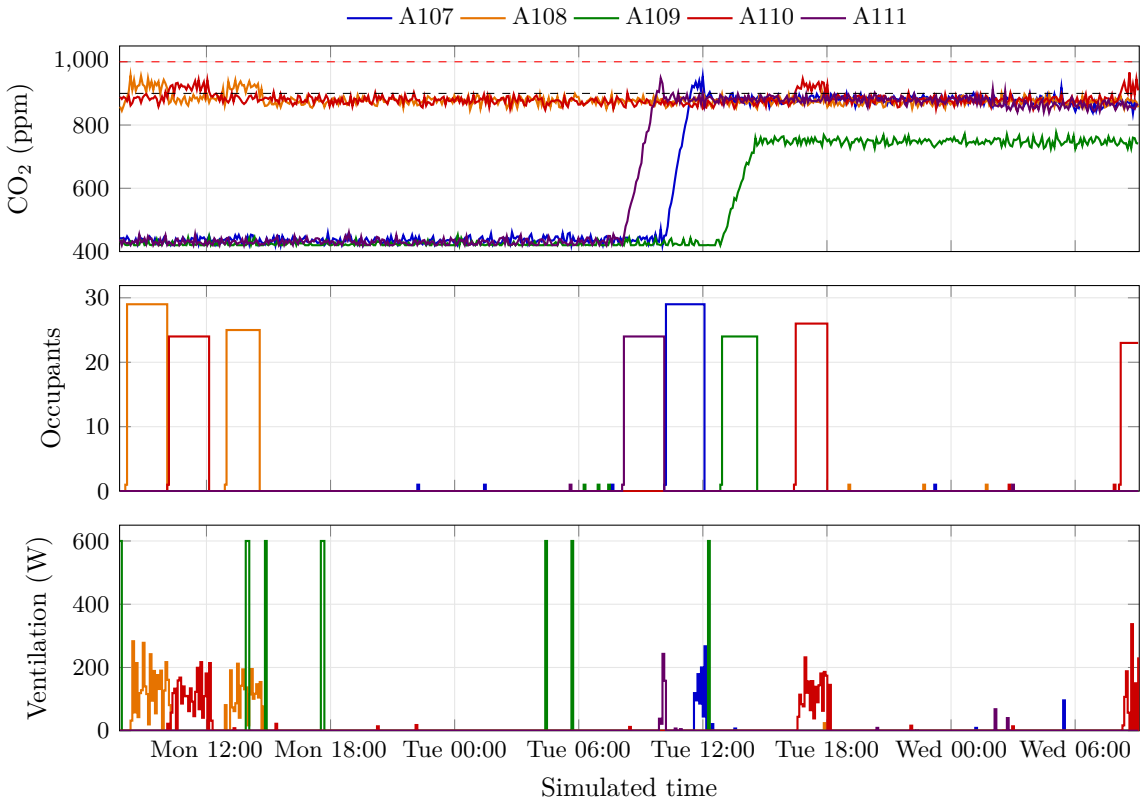


Figure 11: E1: CO_2 (dashed lines at the 900 ppm setpoint and the 1000 ppm limit), head-count from the occupancy simulator, and ventilation command per room.

Five rooms hosted lectures of 24–29 people (Figure 11). CO_2 rose during each lecture and peaked at 949–966 ppm, just above the 900 ppm setpoint, where the ventilation loop held it with up to 337 W (median 109 W). In 118 steps a room's CO_2 was above 900 ppm, and in every one of them the ventilation was on in the next step, which is the one-step actuation lag of Section 6. No occupied step exceeded 1000 ppm, so FR-3 is met as measured, with a caveat that matters: the rooms are four times too large (Section 13), which slows the CO_2 rise by the same factor. With the real volume the replay puts three rooms above 1000 ppm for 7–26 % of their occupied steps.

Two behaviours in the figure are model artefacts, not control behaviour. After a lecture, CO_2 stays where it was, because the CO_2 balance has no infiltration term and the controller only ventilates above 900 ppm; A110 kept about 880 ppm all night. And A109's ventilation shows 600 W spikes with no one present: six false fail-safes, nine steps at full ventilation. A109's true CO_2 sat at 420 ppm for most of the run, the simulator clamps noisy readings at 420 ppm, so 29 % of A109's readings were exactly 420.0, and five identical readings in a row trip the stuck-sensor rule. The data-quality check and the sensor model interact in a way neither unit test could show. FR-4 is met: occupancy reached BuildSim every step, with the lecture blocks visible in all five rooms.

11.3 Cadence and latency (NFR-1)

Over the nominal part of E1 the simulator stepped every 5.08 s (median; maximum 5.67 s) and needed 21 ms of work per step for all five rooms (maximum 84 ms). The decision maker made 5808 of the 5810 expected decisions, two per room per step; the two missing ones are the very first step, before its subscriptions were in place. Measured on the Docker log clock, a reading took 27 ms from the simulator's publish to the decision maker (median; 107 ms at p95), and a command

29 ms from decision to confirmed actuator (100 ms at p95), so the whole loop took 59 ms at the median and 207 ms at the 95th percentile, with a worst case of 886 ms (Figure 12). That is five times inside the NFR-1 budget and 1–4 % of the 5 s step. The tail most likely comes from the sensors opening a new MQTT connection per message (`publish.single`) and from the HTTP round trips to BuildSim inside the sensors and actuators; we did not profile it, because at this step length it does not matter.

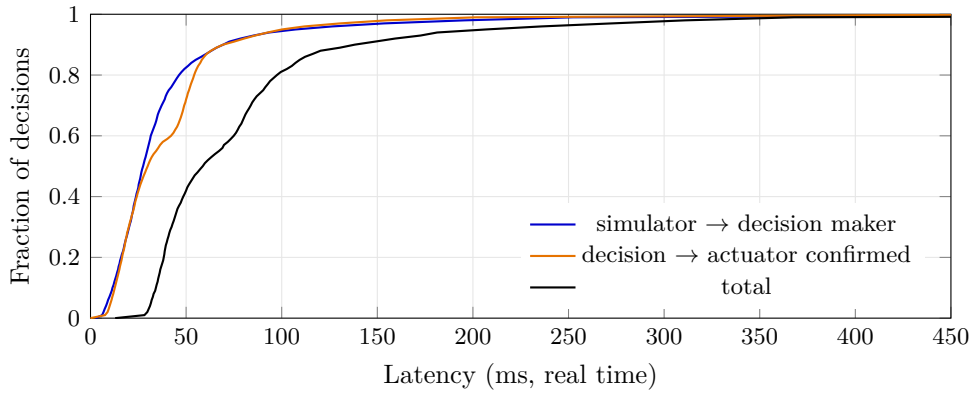


Figure 12: E1: cumulative distribution of per-decision latency over 5808 decisions. The requirement (NFR-1) is 1000 ms at the 95th percentile.

11.4 Fault injection (NFR-2 to NFR-5)

Each fault was injected into the running E1 stack with Docker commands, one at a time, with the other four rooms running normally.

Fault	Detection	Behaviour during fault	Recovery
E2: A110 temperature sensor stopped 65 s	Stale 15.9 s after the last reading (15 s limit + 1 s watchdog tick)	Fail-safe: A110 heating 5.7 kW → 0 W. Other rooms made all 104 expected decisions	PID resumed 0.65 s after the sensor restarted
E3: A109 HVAC actuator stopped 44 s	Resend 10.9 s after the first unconfirmed command; 3 resends	A109 heater held its last applied power; other rooms made all 72 expected decisions	Command confirmed 2.5 s after the actuator restarted
E4: decision maker killed (SIGKILL)	None: actuators cannot tell	All actuators held their last value; no decisions	Not restarted by Docker; after a manual start, first decision 24 s after the kill
E5: MQTT broker stopped 22 s	All 10 loops stale after 15 s	Fail-safe commands undeliverable (actuator resend warnings); simulator crashed 3 times (broker hostname unresolvable) and was restarted	All 10 loops back under control 8.1 s after the broker returned

The sensor and actuator faults behaved exactly as designed, including the isolation: during both outages the other four rooms lost no decisions (NFR-2, NFR-3, NFR-4). The actuator timeout fired after 10.9 s rather than 10 s because the watchdog checks once per second and the timeout is strict; the criterion should have allowed for the tick. The interesting results are E4 and E5. Docker's `unless-stopped` policy restarts a container whose process exits, as the simulator showed three times in E5, but it treats `docker kill` as a deliberate stop, so a killed decision maker stays down; a crash would have been recovered, a kill was not. And in E5 the fail-safe worked on the decision side but could not reach the actuators: with QoS 0 and no broker, the fail-safe commands were simply lost, so the actuators held their last values for the 22 s the broker was gone. That is

the right outcome only by luck; a longer outage in winter would leave heaters running on stale commands. An actuator-side watchdog (revert to a safe value if no command arrives for a while) is the fix, and is in Section 13.

11.5 Scaling and the breaking point (NFR-6)

The same procedure (start the stack, wait 150 s after the first step, measure) was run with 5, 20 and 40 rooms; the five-room run reuses E1's rooms and larger runs add the next rooms that resolve cleanly (Section 13). Memory is what limits the design (Figure 13). Each controlled room adds five Python containers of about 19 MiB each as Docker reports them, but the VM's used memory grows by 116 MiB per room, 23 MiB per container once the runtime's per-container overhead is included. At 40 rooms the stack ran 208 containers in 5966 MiB of the VM's 7789 MiB, with no restarts and no missed decisions; the linear fit reaches the VM's memory at about 56 rooms.

Latency sets a similar limit. The shared simulator's work per step grew from 27 ms to 253 ms, harmless against a 5 s step, but the 95th percentile of the whole loop grew from 123 ms to 704 ms, by about 21 ms per room between 20 and 40 rooms, which would cross NFR-1's 1 s budget near 54 rooms. Both halves of the loop slow down (at the median, sensing from 108 to 319 ms and actuation from 70 to 218 ms between 20 and 40 rooms), which points at per-message work that serialises the burst of 2N messages each step produces: every hop makes an HTTP call to BuildSim, and the decision maker handles messages one at a time, calling the occupancy simulator for every temperature reading to look up the setpoint. We did not profile which dominates; caching the setpoint once per step is the cheap first fix.

The projection for the earlier design (eight containers per room, no shared services) shows why consolidation was necessary but not sufficient. It would fill the VM at about 35 rooms, and its 952 containers for all 119 configured rooms would need about 23 GB, so it could not have run on this laptop even with the occupancy simulator's health check fixed. The current design runs 40 rooms comfortably but would still need about 11 GB for the 84 rooms that resolve cleanly: one container per emulated device is right for demonstrating fault isolation and wrong for a whole building.

The 40-room run also found an operational limit. With 202 services sharing one build section, `docker compose up -build` failed inside BuildKit ("no such job"); building the shared image once and then starting with `up -d` works, and `run_scale.sh` now does that.

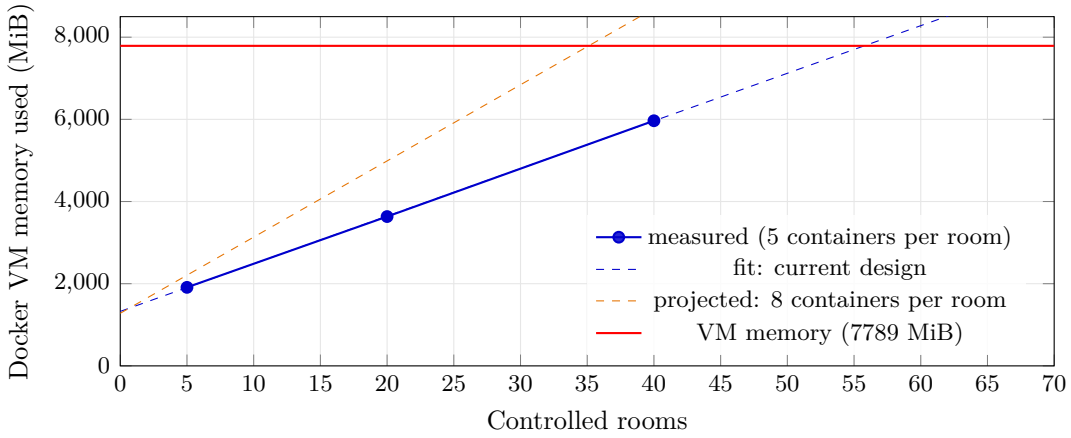


Figure 13: Memory used in the Docker VM against the number of controlled rooms, measured 150 s after the first simulation step, with a linear fit and the projection for the earlier design with per-room simulator and decision makers.

Per-run measurements. Each row is one run of `run_scale.sh`, measured 150 s (about 30

steps) after the first simulation step; the decisions column excludes each run's first and last step.

Rooms	Containers running	VM memory used (MiB)	Simulator work per step, p50 / max (ms)	Latency p50 / p95 (ms)	Decisions missed / expected	Restarts
5	33	1911	27 / 71	63 / 123	0 / 290	0
20	108	3634	111 / 166	193 / 280	4 / 1160	0
40	208	5966	253 / 429	571 / 704	0 / 2320	0

11.6 Results against the requirements

Req.	Criterion	Measured	Verdict
FR-1	Reach 21 °C in 24 h, then ±1 °C (T-1); live: ≥ 90 % of daytime steps in band	T-1 passes. Live: 35–63 % of daytime and 52 % of occupied steps in band; mornings need 2.4–6.3 h	× live
FR-2	Switch within one step; less energy than constant 21 °C	All four switches in the first step; 30 % less HVAC energy (replay), at a large comfort cost	✓
FR-3	Ventilation on above 900 ppm; full room below 1000 ppm in 10 min	On in the next step in 118 of 118 cases; no occupied step above 1000 ppm (peak 966); T-2 passes	✓[c]
FR-4	Head-count in BuildSim every step	Every step; lectures of 24–29 people in all rooms; T-3 passes	✓
FR-5	Rooms chosen by flag only	5, 20 and 40 rooms deployed by flag and generator; T-7	✓
NFR-1	Every room every step; p95 < 1 s	5808 of 5810 decisions; p95 207 ms at 5 rooms, 704 ms at 40	✓
NFR-2	Fail-safe within 1 s of the 15 s limit; resume on good data	15.9 s; resumed 0.65 s after the sensor returned; but six false stuck alarms in E1	✓[c]
NFR-3	Detect within 10 s; resend until confirmed	10.9 s (1 s watchdog tick); 3 resends; confirmed 2.5 s after restart	≈
NFR-4	Other rooms unaffected	All 104 and 72 expected decisions made during E2 and E3; T-6	✓
NFR-5	Control back within 60 s, no manual action	Broker outage: 8.1 s. Killed decision maker: not restarted by Docker	×
NFR-6	5 rooms < 2 GB; ≥ 20 rooms in 8 GB	5 rooms: 1911 MiB in the VM; 40 rooms: 5966 MiB; limit about 55 rooms	✓

[c] Met as measured, with the caveat given in the text (FR-3: rooms four times too large; NFR-2: false alarms).

Eight of the eleven requirements are met as measured, two of them with a caveat; NFR-3 is missed by the watchdog's one-second tick; FR-1 fails in the live system and NFR-5 fails for a killed decision maker. The failures share a pattern worth stating plainly: every component passed its unit test, and every failure appeared only when the whole system ran for long enough, with real occupancy, a real schedule and real faults. That is the strongest argument this project produced for end-to-end evaluation, and the reason Section 13 is longer than we expected.

12. Dashboard

The dashboard is a Grafana board provisioned from `monitoring/` when the stack starts, so it needs no manual setup and lives in version control. It shows the outside temperature, the occupancy simulator's clock in Swedish time, and six per-room time series: temperature, setpoint, HVAC power, CO_2, ventilation command and occupancy. A room selector filters the room panels, and the board refreshes every 5 s real, once per simulated step. BuildSim's own 3D viewer, open next

to it, shows the same sensor values on the rooms and the occupancy simulator's people moving; the two views share a source of truth, so they cannot disagree.

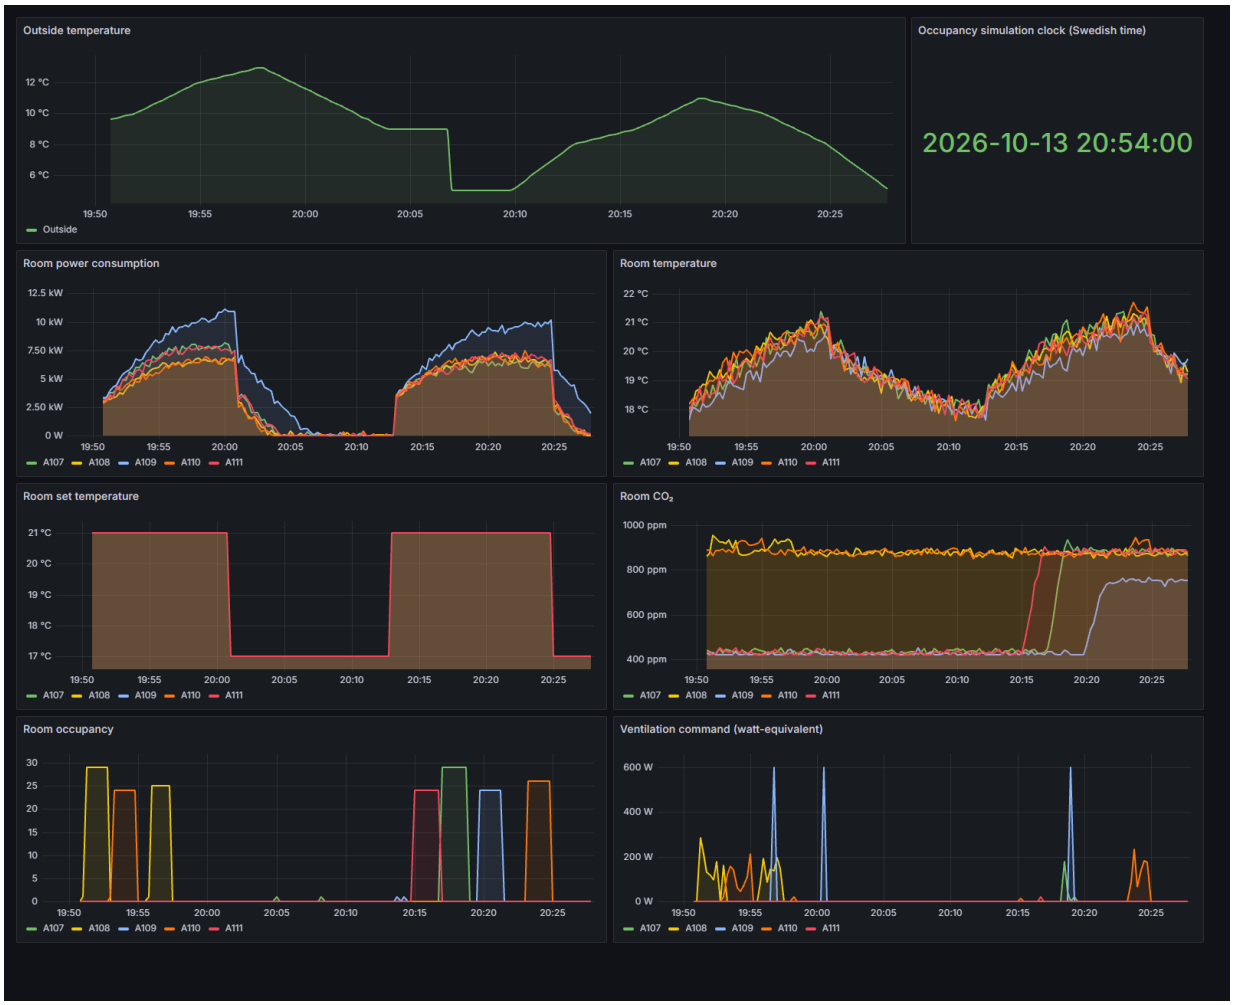


Figure 14: The dashboard at the end of E1's nominal window, captured headless from the running stack. The x-axis is real time (37 real minutes, about 37 simulated hours); the clock panel shows the simulated time.

The screenshot tells the E1 story at a glance: the setpoint panel shows the 21/17 °C schedule, the power panel the heating ramping up on two mornings, halving at 18:00 and winding down over the evening, the occupancy panel the lectures arriving in different rooms each day, and the CO_2 panel rooms rising during lectures and then staying at the new level. The ventilation panel also shows something we did not intend: 600 W spikes in A109 with no one in the room. They are false fail-safes, explained in Section 11.4; the dashboard is how we found them.

What the dashboard does not show is the decision maker's internal state: whether a loop is in fail-safe or waiting for an acknowledgement is visible only in the container log. That is the first thing we would add (Section 13).

13. Risks and critical reflection

Writing the evaluation found more problems than building the system did, and the most useful ones are below, worst first. None of them was visible from the unit tests alone; they came from running the whole stack and reading the numbers.

1. The rooms are four times too large. BuildSim reports room `area` in plan units² (0.25 m² each at the occupancy simulator's 0.5 m per unit), and `con-sim.py` uses that number as m². A107 is 122 m² to the occupancy simulator but 490 "m²" to our physics, so every room's air volume, thermal mass and envelope are four times too large. This one bug explains two results that looked like design properties: the slow warm-up (Section 11.1) and the comfortable CO_2 margin (Section 11.2). Replaying E1 with the area corrected, the same controller reaches 20.5 °C in 0.7–3.8 h instead of 5.4–10.1 h, but A108's CO_2 peaks at 1019 ppm and is above 1000 ppm for 15 % (A107: 26 %) of its occupied steps: the ventilation gains that look adequate today are too weak for a real lecture room. Fixing the conversion is one line; retuning the CO_2 loop is the follow-up it forces.

2. The simulator stops every simulated weekend. `Weather.json` covers Monday to Friday only, and the outside-temperature lookup sits outside the per-room error handling, so on Saturday the simulator raises, Docker restarts it, and it raises again at the first step. At factor 60 that happens 1.9 real hours after a Monday 07:30 start and lasts 48 real minutes. The decision maker then correctly fails safe in every room, so the building is not harmed, but no physics runs. Either the data or a fallback (repeat Friday) is needed before a long demo.

3. The controller integrates real time, the plant simulated time. The PID uses `time.monotonic()`, so K_i and K_d are effectively per real second, while the plant evolves 60 times faster. The gains are therefore only valid at factor 60: at factor 120 the integral action per simulated hour halves. Readings should carry the simulated timestamp, and the PID should use it. A related detail showed up at every 18:00: the integral is not reset when the setpoint drops, so the heating takes 2–6 hours to wind down instead of stopping.

4. The schedule ignores both occupancy and thermal lag. The 06:00 setpoint step reaches the rooms only after 2.4–6.3 hours (Section 11.1), so lectures at 08:15 start below the comfort band, while four of the five rooms sat empty for most of the day and were heated anyway. The setback saved 30 % of the HVAC energy against a constant 21 °C; an occupancy-aware setback with an optimal-start ramp (start heating early enough to reach 21 °C when the first lecture begins) would address both the comfort miss and the remaining waste, and is the main improvement to the autonomous part.

5. CO_2 never decays without mechanical ventilation. The CO_2 balance has no infiltration term, although the heat balance does ($ACH_0 = 0.5\ h^{-1}$). An empty room therefore keeps its last CO_2 value until the damper opens, which is why A110 sat at about 870 ppm overnight. Adding the same background air change to the CO_2 balance fixes it.

6. The stuck-sensor rule raises false alarms. Five identical readings mean "stuck" only if a live sensor never repeats itself. Our simulated CO_2 sensor does repeat itself whenever the true value is at the 420 ppm floor, because noise below the floor is clamped, and A109 failed safe six times with nobody in the room (Section 11.2). It is systematic, not a one-off: in the 20-room scaling run, newly enabled rooms start at the 420 ppm default and seven of the fifteen raised the same false alarm within 2.5 minutes. The cost per alarm is small (a step or two of full ventilation), but the lesson is general: a data-quality rule encodes an assumption about the sensor, and here the sensor model broke it. The rule should ignore readings at a physical bound, or the simulator should not clamp.

7. The shared services are single points of failure. Consolidation (D2, D3) traded per-room isolation for a system that starts. We mitigated rather than removed the risk: in-process isolation, a restart policy and the fail-safe. E4 and E5 show what is left: killing the decision maker left every actuator at its last value, and Docker did not restart a container it had killed itself, so control came back only after a manual start (24 s after the kill, 20 s of which was our wait). The broker outage needed no help: all loops resumed 8 s after the broker returned, although the simulator crashed and restarted three times while it was gone. The same consolidation has a subtler cost: the shared simulator refuses to start if any enabled room is missing from BuildSim and waits

indefinitely for any room whose devices never register, so one misconfigured room now stops the physics for all of them. And 20 of the 119 rooms in `rooms.json` have no unique match in the occupancy simulator; enabling all rooms today would crash those rooms' occupancy publishers in a loop. The generator should validate rooms before writing the compose file.

8. The fail-safe needs the broker it is protecting against. The decision maker decides to fail safe, but the actuators only hear it over MQTT. In E5 the broker was the failure, so the fail-safe commands were lost and the actuators held their last values for 22 s. A watchdog in each actuator (revert to a safe value when no command has arrived for, say, three steps) would make the fail-safe independent of the network, which is where a real BMS puts it.

9. Per-device containers cost memory. Each Python container takes about 23 MiB of VM memory including Docker's overhead, so five devices per room dominate the footprint. The measured breaking point (Section 11.5) is about 55 rooms on this laptop, where the 8 GB VM fills up and the loop latency reaches its 1 s budget at roughly the same point. The device-per-container principle is right for a demonstration of fault isolation; a real deployment would run many virtual devices per process, or real devices.

10. The fail-safe for heating is "off". That is safe against a frozen reading driving the heater, but a long sensor outage in winter cools the room without bound. A frost-protection minimum instead of 0 W would be the better default for a real building.

11. Monitoring has no memory and no view of decisions. Prometheus stores to tmpfs, so a restart erases the evaluation data unless it is exported first (we exported ours), and the dashboard cannot show whether a loop is in fail-safe or retrying. Exporting the controller state as metrics would be a small change with a large effect on trust.

12. No authentication anywhere. MQTT accepts anonymous clients and BuildSim has no authentication. Publishing every port only on 127.0.0.1 keeps this acceptable on a laptop, but any process on the machine can command an actuator.

Risk / assumption	Impact	Mitigation / next step
Area in plan units used as m^2	Wrong dynamics; CO_2 loop under-tuned	Convert $\times 0.25$; retune ventilation gains
No weekend weather data	No physics 48 h per simulated week	Add data or repeat Friday; catch the error per step
PID on real time	Gains tied to factor 60	Use simulated timestamps in readings
Clock-only setback	Cold mornings, heated empty rooms	Occupancy-aware setback with optimal start
Stuck rule meets clamped sensor	False fail-safes (6 in E1)	Ignore readings at a physical bound
Shared `dec/sim`	All rooms affected by one crash; `kill` not restarted	Restart policy covers crashes (E5); external health check for the rest
Fail-safe travels over MQTT	Lost when the broker is down (E5)	Actuator-side watchdog; QoS 1 for commands
Five containers per room	Memory caps the room count	Several devices per process beyond a demo
Unvalidated `rooms.json`	Crash loops for 20 of 119 rooms	Validate against BuildSim and occupancy simulator in the generator

What we would do next, in order: fix the area conversion and retune the CO_2 loop; make the simulator survive weekends and broker outages; feed simulated time to the PID and reset its integral on setpoint changes; stop clamping the simulated CO_2 sensor; add a watchdog in each actuator; replace the clock setback by an occupancy-aware optimal start; export controller state to the dashboard. The first four are bugs that the evaluation found; the last three are the design work that would make the autonomous part genuinely better.